

# **Operating Systems**

## **Complex Computing Problem**

---



**Muhammad Shakaib Arsalan (F2022266626)**

**Course Code: CC3011L**

**Section: V5**

**Course Instructor: Habiba Habib**

School of Systems and Technology

UMT Lahore Pakistan

# Estimating Pi using the Maclaurin Series

We have to calculate the value of  $\pi$  using the Maclaurin series for  $\tan^{-1}(x)$ . This is the mathematical computational problem. Find  $\pi$  using Maclaurin series for is one the effective approach.

We can express  $\pi$  as:

$$\pi = 4 * \sum_{n=0}^{\infty} \frac{-1^n}{2n+1}$$

$$\pi = 4 \times \left(1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \frac{1}{9} - \dots\right) \quad n = 1$$

Eq shows an optimized version of the Maclaurin Series for  $\tan^{-1}(1)$  and  $\tan^{-1}(1) = \frac{\pi}{4}$

The goal is to implement parallel computing approach to approximate  $\pi$  using series. By ensuring the equal distribution of workload using multithread and multiprocessing.

## 1. Sequential Implementation

### Algorithm:

The sequential implementation uses a single process and thread, and estimates PI in a single loop.

### Code:

Click to access [sequential.cpp](#) for the complete implementation.

### Steps:

1. Read the number of terms,  $n$ , from the command line.
2. Initialize a global variable PI to store the computed result.
3. Loop from 0 to  $n - 1$  and compute the sum using the formula.
4. Multiply the result by 4 to get the final value of  $\pi$ .

### Discussion:

This solution is simple but lacks parallelism, resulting in high execution time for larger values of  $n$ . It does not efficiently utilize available system resources.

## 2. Multiprocessor Implementation Details

### Algorithm:

The multiprocessor implementation uses shared memory with processes to distribute computation.

### Steps:

Create shared memory for storing the global PI result.

1. Fork child processes based on the number of requested processes.
2. Each process computes a partial sum of the Maclaurin series and writes to shared memory.
3. Use locks to synchronize access to the shared  $\pi$  variable.
4. Parent process waits for all child processes to complete.

### Code:

Click to access [shared-memory.cpp](#) for the complete implementation.

### Discussion:

This solution achieves better performance than the sequential implementation by utilizing multiple processes. However, synchronization through shared memory and locks introduces some complexity and potential contention.

## 3. Multithreading-based Solution

### Algorithm:

The multithreading approach divides the computation across multiple threads, synchronizing access to a shared  $\pi$  variable using mutex locks.

### Steps:

1. Initialize a global variable  $\pi$  and a mutex lock.
2. Spawn threads based on the user-specified thread count.
3. Each thread computes a partial sum of the Maclaurin series.
4. Threads acquire the lock to update the shared  $\pi$  variable and release it afterward.
5. Wait for all threads to finish computation.
6. Compute the final value by multiplying  $\pi$  by 4.

**Code:** Click to access [thread.cpp](#) for the complete implementation.

### Discussion:

This solution provides faster execution and better resource utilization compared to the process-based approach. Proper lock management ensures accurate results.

## **4. Results and Performance Analysis**

### **Test Environment:**

Processor: Intel Core i7 5<sup>th</sup> gen.

Operating System: Kali-Linus-2024.3

### **Performance Comparison:**

Metric	Sequential	Multiprocessing	Multithreading
Execution Time	Slow	Moderate	Fast
Sync. Overhead	None	High	Low
Resource Utilization	Low	High	High
Scalability	Poor	Moderate	Excellent

### [Performance Comparison Graph](#)

### **Observations:**

1. The multithreading solution provided the best performance due to reduced context-switching overhead compared to processes.
2. Synchronization was handled efficiently using mutex locks.
3. The sequential approach was significantly slower, demonstrating the necessity for parallel solutions.

## **5. Discussion on Locks and Synchronization**

Mutex locks are used to enter in critical sections when updating the global PI variable.

exact placement of locks ensured correctness of value of updated PI

Shared memory with atomic operations was used in the process-based solution to synchronize access.

## **6. Code Files**

1.  [sequential.cpp](#): Sequential Implementation
2.  [shared-memory.cpp](#): Multiprocessor Implementation using Shared Memory
3.  [pipes.cpp](#): Process Communication using Pipes
4.  [thread.cpp](#): Multithreading-based Solution

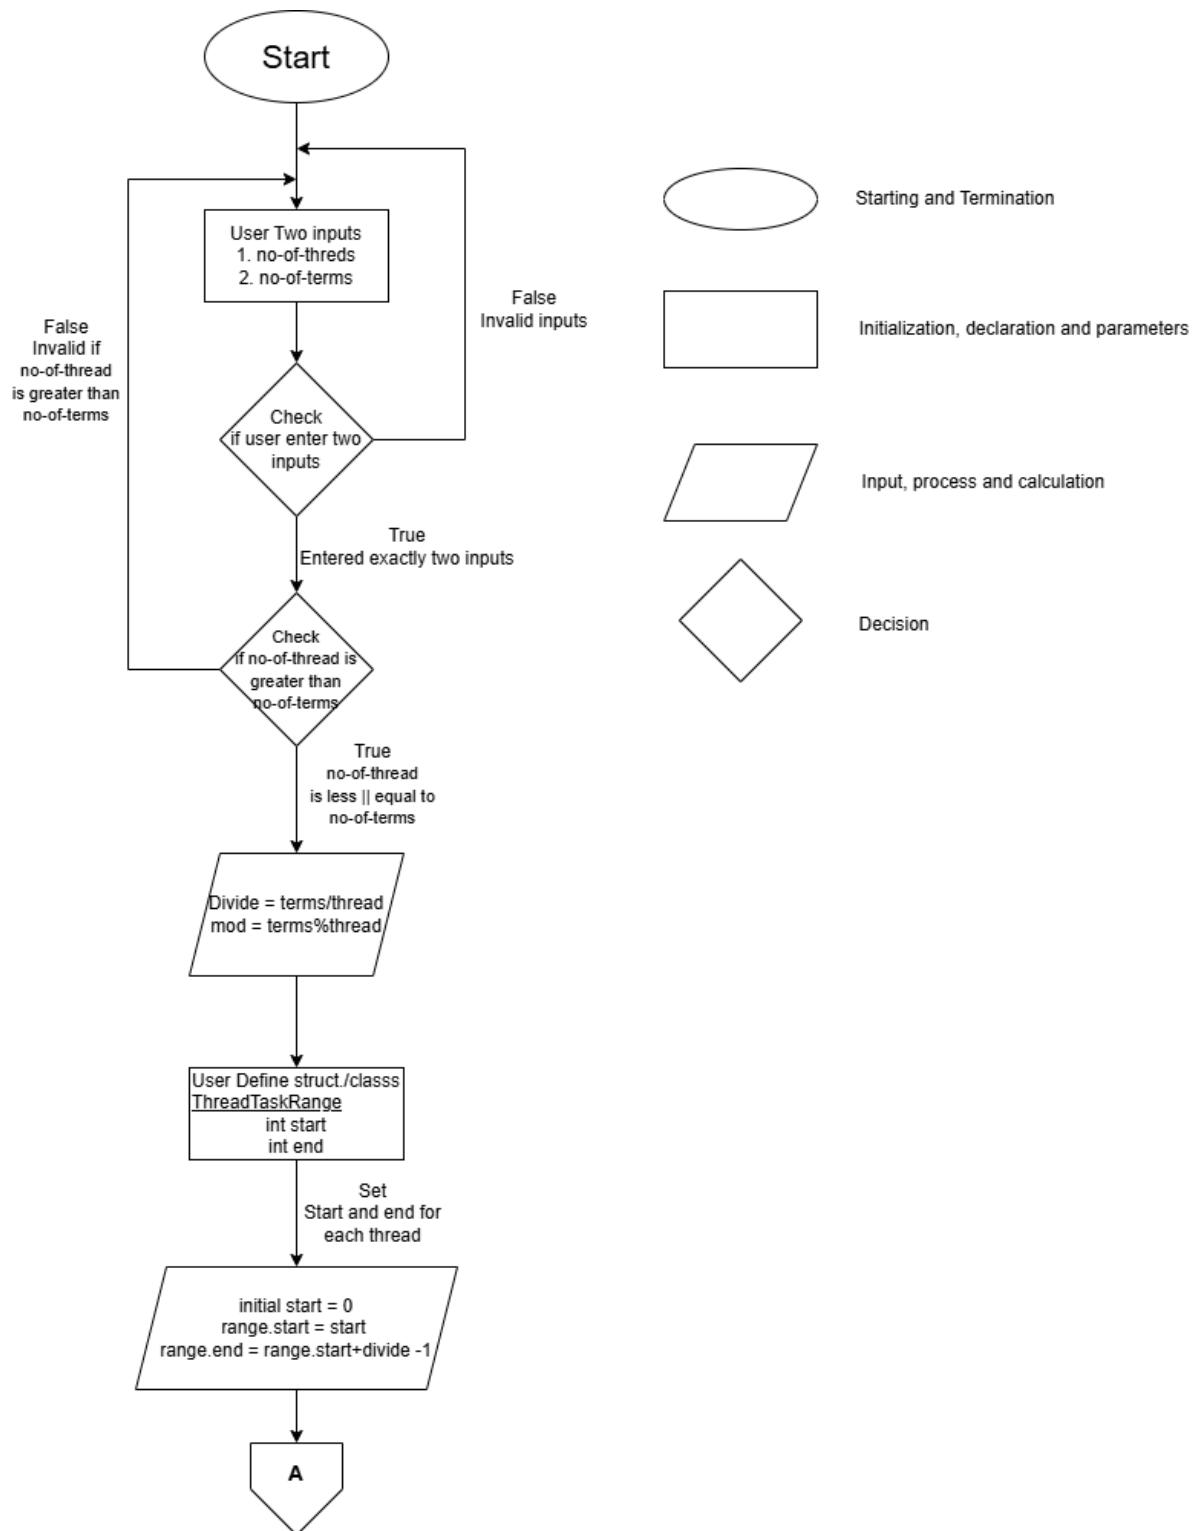
## **7. Performance Analysis Discussion**

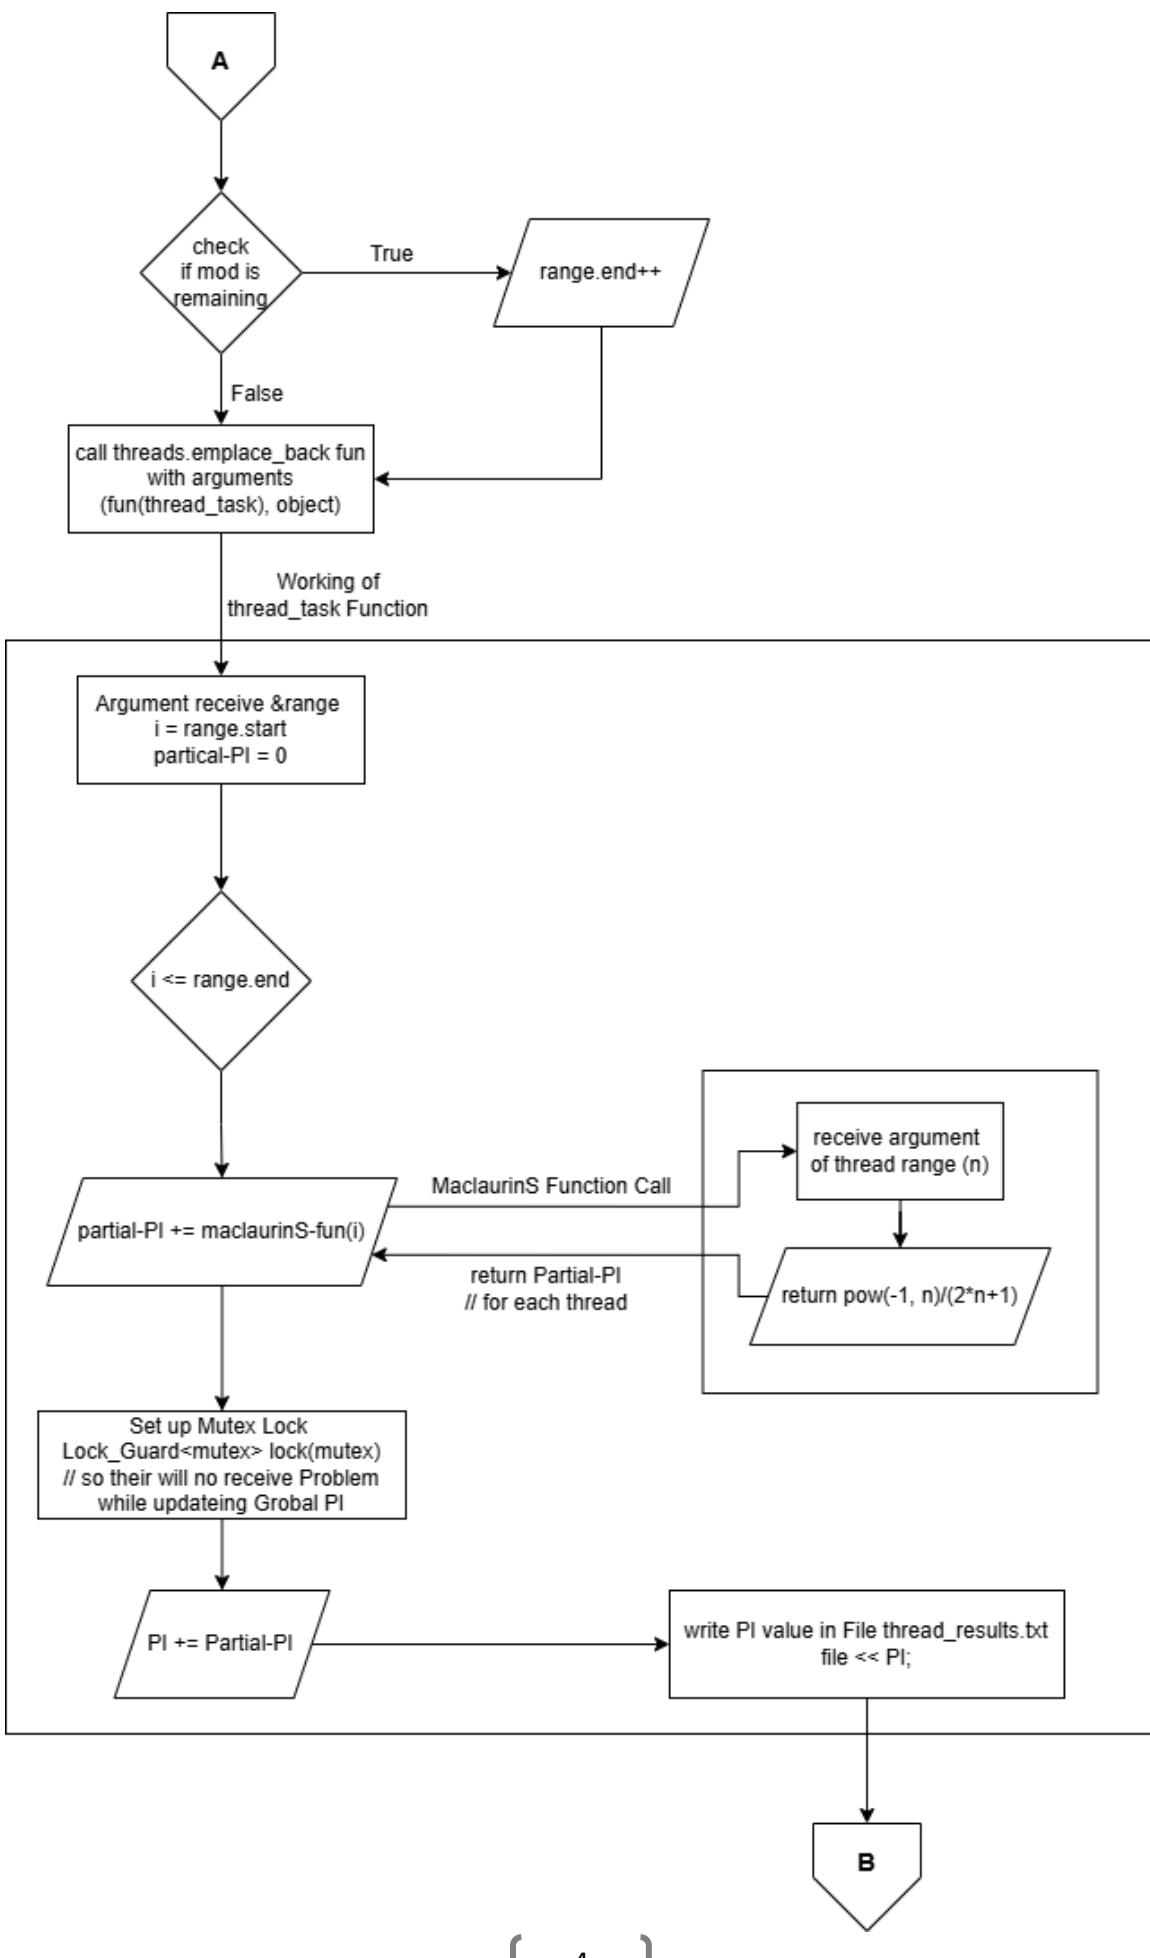
1. The parallel solutions showed significant improvement over the sequential implementation.
2. Proper workload distribution and synchronization played a crucial role in achieving accurate and efficient results.
3. Load balancing was achieved by evenly dividing the terms among threads and processes.

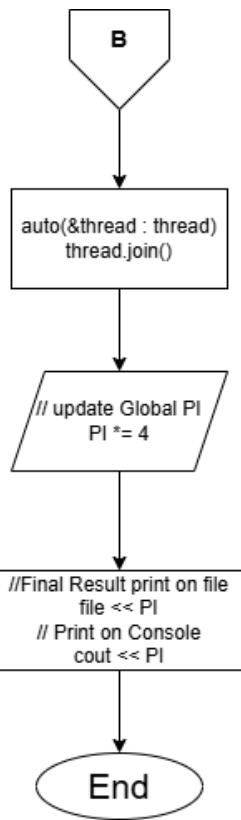
## **8. Conclusion**

This project successfully demonstrated the computation of Pi using the Maclaurin series. Different parallel approaches were implemented and compared, highlighting the advantages of multithreading for computationally intensive tasks. Effective lock management and synchronization ensured accuracy and efficiency.

# Flowchart







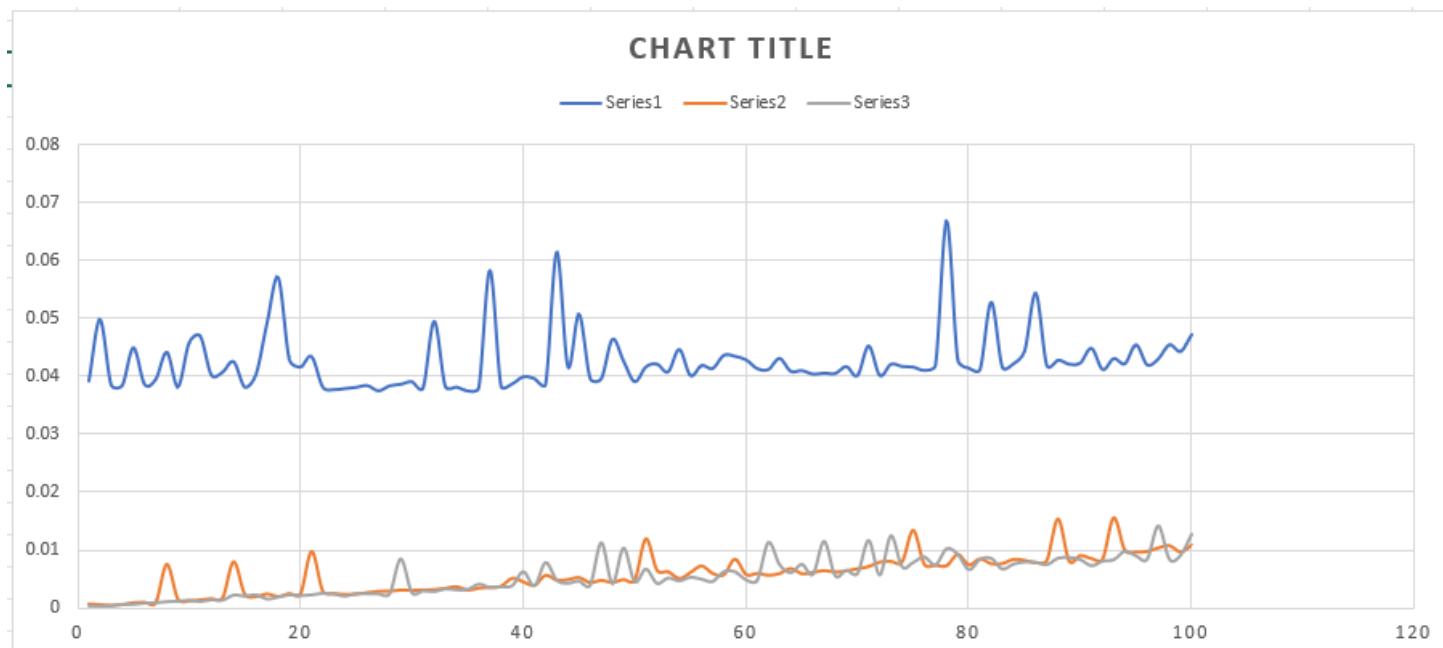
## Performance Comparison:

X is instances from 1 to 100

**Series 1:** Multithreaded

**Series 2:** Shared Memory

**Series 3:** Pipes



## Bash Script:

The screenshot shows a terminal window with two tabs: 'script.sh' and 'multithreaded.c'. The current tab is 'script.sh'. The command history shows the user navigating through a directory structure and then running the script. The script itself is a bash script that loops from 1 to 100, executing three programs ('./b', './c', and './d') for each iteration with a parameter of 5000000.

```
script.sh    ●  C multithreaded.c ●
: Sem III > OS > CCP > codes (copy) > $ script.sh
1  #!/bin/bash
2
3  # Loop from 1 to 100
4  for i in {1..100}; do
5      ./b "$i" 5000000
6      ./c "$i" 5000000
7      ./d "$i" 5000000
8  done
9
10 |
```

## Code

### sequential.cpp

```
1
2  #include <iostream>
3  #include <cstdlib>
4  #include <cmath>
5  using namespace std;
6
7  long double maclaurin_polynomial_for_PI(unsigned int n)
8  {
9      return pow(-1, n) / (2 * n + 1);
10 }
11
12 double compute_pi(unsigned int n)
13 {
14     double pi = 0;
15     for (unsigned int i = 0; i < n; i++)
16     {
17         pi += maclaurin_polynomial_for_PI(i);
18     }
19     return pi * 4;
20 }
21
22 int main(int argc, char *argv[])
23 {
24     if (argc != 2)
25     {
26         cout << "Invalid argument count\n";
27         return 1;
28     }
29
30     unsigned int n = static_cast<unsigned int>(atoi(argv[1]));
31     double PI = compute_pi(n);
32
33     cout << "PI: " << PI << endl;
34
35     return 0;
36 }
```

## pipes.cpp

```
1 #include <iostream>
2 #include <cstdlib>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <cmath>
6
7 using namespace std;
8
9 long double maclaurin_polynomial_for_PI(unsigned int n) {
10     return ((n % 2) ? -1.0L : 1.0L) / (2 * n + 1);
11 }
12
13 long double child_func(int start, int end) {
14     long double local_sum = 0.0L;
15     for (unsigned int i = start; i <= end; i++) {
16         local_sum += maclaurin_polynomial_for_PI(i);
17     }
18     return local_sum;
19 }
20
21 int main(int argc, char *argv[]) {
22     if (argc != 3) {
23         cerr << "Usage: " << argv[0] << " <num_processes> <num_terms>" << endl;
24         return 1;
25     }
26
27     int no_of_processes = atoi(argv[1]);
28     int no_of_terms = atoi(argv[2]);
29
30     if (no_of_processes < 1 || no_of_terms < 1) {
31         cerr << "Number of processes and terms must be at least 1." << endl;
32         return 1;
33     }
34
35     int pipes[no_of_processes][2];
36     int div_each = no_of_terms / no_of_processes;
37     int mod_remain = no_of_terms % no_of_processes;
38     int start = 0;
39
40     for (int i = 0; i < no_of_processes; i++) {
41         if (pipe(pipes[i]) == -1) {
42             perror("Pipe error");
43             return 1;
44         }
```

```

46     int end = start + div_each - 1;
47     if (i < mod_remain) end++; // Distribute remainder
48
49     pid_t pid = fork();
50
51     if (pid == -1) {
52         perror("Fork error");
53         return 1;
54     }
55     else if (pid == 0) { // Child process
56         close(pipes[i][0]); // Close read end
57         long double partial_sum = child_func(start, end);
58         write(pipes[i][1], &partial_sum, sizeof(long double));
59         close(pipes[i][1]);
60         exit(0);
61     }
62
63     start = end + 1; // Update for next process
64 }
65
66 long double PI = 0.0L;
67
68 // Parent process collects results
69 for (int i = 0; i < no_of_processes; i++) {
70     close(pipes[i][1]); // Close write end
71     long double partial_sum;
72     read(pipes[i][0], &partial_sum, sizeof(long double));
73     close(pipes[i][0]);
74     wait(NULL); // Wait for child
75     PI += partial_sum;
76 }
77
78 PI *= 4;
79 cout.precision(20);
80 cout << "\nPI is: " << PI << endl;
81
82 return 0;
83 }
```

## shared-memory.cpp

```
1 #include <iostream>
2 #include <cstdlib>
3 #include <unistd.h>
4 #include <sys/wait.h>
5 #include <sys/mman.h>
6 #include <cmath>
7 #include <atomic>
8 using namespace std;
9
10 typedef atomic_flag lock_t;
11
12 // #define LOCK_INIT ATOMIC_FLAG_INIT
13
14 static inline void acquire(lock_t *lock)
15 {
16     while (lock->test_and_set(memory_order_acquire));
17 }
18
19 static inline void release(lock_t *lock)
20 {
21     lock->clear(memory_order_release);
22 }
23
24 long double maclaurin_polynomial_for_PI(unsigned int n)
25 {
26     return ((n % 2) ? -1.0L : 1.0L) / (2 * n + 1);
27 }
28
29 void child_func(int start, int end, long double *PI, lock_t *mutex)
30 {
31     long double local_sum = 0.0L;
32     for (int i = start; i <= end; i++)
33     {
34         local_sum += maclaurin_polynomial_for_PI(i);
35     }
36
37     acquire(mutex);
38     *PI += local_sum; // Write to shared memory
39     release(mutex);
40 }
41
42 int main(int argc, char *argv[])
43 {
44     if (argc != 3)
45     {
46         cerr << "Invalid argument count" << endl;
47         return 1;
48     }
49
50     int no_of_child = atoi(argv[1]);
51     int no_of_terms = atoi(argv[2]);
52
53     if (no_of_child < 1 || no_of_terms < 1)
54     {
55         cerr << "Number of processes and terms should be at least 1." << endl;
56         return 1;
57     }
```

```

60     long double *PI = (long double *)mmap(NULL, sizeof(long double), PROT_READ | PROT_WRITE, MAP_SHARED | MAP_ANONYMOUS, -1, 0);
61     lock_t *mutex = (lock_t *)mmap(NULL, sizeof(lock_t), PROT_READ | PROT_WRITE, MAP_SHARED | MAP_ANONYMOUS, -1, 0);
62
63     if (PI == MAP_FAILED || mutex == MAP_FAILED)
64     {
65         perror("mmap error");
66         return 1;
67     }
68
69     *PI = 0.0L;
70     // *mutex = LOCK_INIT;
71     atomic_flag_clear(mutex);
72
73     int div_each = no_of_terms / no_of_child;
74     int mod_remain = no_of_terms % no_of_child;
75     int start = 0;
76
77     for (int i = 0; i < no_of_child; i++)
78     {
79         int end = start + div_each - 1;
80         if (i < mod_remain)
81             end++;
82
83         pid_t pid = fork();
84         if (pid == -1)
85         {
86             perror("fork error");
87             return 1;
88         }
89         else if (pid == 0)
90         { // Child process
91             child_func(start, end, PI, mutex);
92             exit(0);
93         }
94         start = end + 1;
95     }
96
97     // Parent waits for all children to finish
98     for (int i = 0; i < no_of_child; i++)
99     {
100         wait(NULL);
101     }
102
103     *PI *= 4;
104     long double error = 100 - ((*PI / 3.1415926535) * 100);
105     cout.precision(20);
106     cout << "\n\nPI is: " << fixed << *PI << endl;
107     cout << "error: " << error << endl;
108
109     munmap(PI, sizeof(long double));
110     munmap(mutex, sizeof(lock_t));
111
112     return 0;
113 }
```

## thread.cpp

```
1 #include <iostream>
2 #include <cmath>
3 #include <vector>
4 #include <thread>
5 #include <mutex>
6 #include <fstream>
7
8 long double PI = 0.0;
9 std::mutex mutex;
10 std::mutex file_mutex;
11
12 double maclaurin_polynomial_for_PI(unsigned int n)
13 {
14     return std::pow(-1, n) / (2 * n + 1);
15 }
16
17 class ThreadTaskRange
18 {
19 public:
20     unsigned int start;
21     unsigned int end;
22 };
23
24 void thread_task(const ThreadTaskRange &range, std::ofstream &file)
25 {
26     long double partialPI = 0.0;
27
28     for (unsigned int i = range.start; i <= range.end; ++i)
29     {
30         partialPI += maclaurin_polynomial_for_PI(i);
31     }
32
33     {
34         std::lock_guard<std::mutex> lock(mutex);
35         PI += partialPI;
36     }
37
38     {
39         std::lock_guard<std::mutex> file_lock(file_mutex);
40         file << "Thread Range: " << range.start << " to " << range.end << ", Partial PI: " << partialPI
41     }
42 }
43
44 int main(int argc, char *argv[])
45 {
46     if (argc != 3)
47     {
48         std::cout << "Invalid argument count\n";
49         return 1;
50     }
51
52     unsigned int no_of_threads = std::stoul(argv[1]);
53     unsigned int no_of_terms = std::stoul(argv[2]);
54
55     std::cout << "Thread count: " << no_of_threads << "\n";
56     std::cout << "# of terms: " << no_of_terms << "\n";
```

```

58     if (no_of_threads > no_of_terms || no_of_terms < 1 || no_of_threads < 1)
59     {
60         std::cout << "Invalid inputs. Ensure no_of_terms >= no_of_threads and both are > 0.\n";
61         return 1;
62     }
63
64     unsigned int div_each = no_of_terms / no_of_threads;
65     unsigned int mod_remain = no_of_terms % no_of_threads;
66
67     unsigned int start = 0;
68
69     std::vector<std::thread> threads;
70     std::ofstream file("thread_results.txt");
71
72     if (!file.is_open())
73     {
74         std::cout << "Error opening file\n";
75         return 1;
76     }
77
78     {
79         std::lock_guard<std::mutex> file_lock(file_mutex);
80         file << "Calculating PI using threads:\n";
81     }
82
83     for (unsigned int i = 0; i < no_of_threads; ++i)
84     {
85         ThreadTaskRange range;
86         range.start = start;
87         range.end = range.start + div_each - 1;
88
89         if (i < mod_remain)
90         {
91             range.end++;
92         }
93
94         threads.emplace_back([=, &file]()
95             | | | | { thread_task(range, file); });
96
97         start = range.end + 1;
98     }
99
100    for (auto &thread : threads)
101    {
102        thread.join();
103    }
104
105    {
106        std::lock_guard<std::mutex> file_lock(file_mutex);
107        file << "\nValue of PI before *4: " << PI;
108    }
109
110    PI *= 4;
111
112    {
113        std::lock_guard<std::mutex> file_lock(file_mutex);
114        file << "\nFinal PI value: " << PI << "\n";
115    }
116
117    file.close();
118    std::cout << "\nPI is: " << PI << "\n";
119
120    return 0;
121 }
```

# Turnitin Plagiarism Report



Page 2 of 15 - AI Writing Overview

Submission ID trn:oid:::1:3145659810

## 0% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely AI-paraphrased.

**Caution: Review required.**

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.

## Detection Groups

### 1 AI-generated only 0%

Likely AI-generated text from a large-language model.

### 2 AI-generated text that was AI-paraphrased 0%

Likely AI-generated text that was likely revised using an AI-paraphrase tool or word spinner.